# Generics: Inference & Accommodation

Greg Restall [1] · Philosophy Department, The University of St Andrews

SEPTEMBER 2024

Social hierarchy and domination are enforced in many ways. Oppression is maintained not only by the force of arms and institutions of state and market, but also by our very habits of *thought* and *talk*. We do not learn languages to describe the world in a disinterested and neutral way: we inherit conceptual schemes that present some things as normal and others as unexpected and different. When we learn that *birds fly*, that *women are emotional*, or that *Chinese people are industrious*, we are taught to expect certain things of birds, women, and people from China. In a literal sense, these judgements form a constellation of *prejudices*—pre-judgements that we bring to any encounter with the world, and which inform our thinking, our planning and our action.

These stereotypes are expressed in *generic* judgements—the kinds of judgements with the form *Fs are G*. We learn to think in generic terms very early in our cognitive development (see Gelman 2003, especially Chapter 8; Gelman, et. al. 2015; Leslie 2008; Kirkpatrick 2024). We teach our children to understand the world, passing on truisms like "*birds fly*", "*cows moo*" and "*dogs bark*". As a child learns how to describe the world around them, they learn what to expect as they explore that world. We learn such judgements early in life, and they are not a crutch we discard as we grow up. Generics remain pervasive in adult thought and talk.

These generics differ from explicitly quantified judgements. To say that *women are emotional* does not go so far as to say that *all* women are emotional (there are counterexamples), and they say more than that merely *some* women are emotional. Generic judgements categorise a class, but they classify that class in a defeasible way. Encountering non-emotional women no more defeats the generic generalisation than the existence of emus counts as evidence against the thought that birds fly.

---

[1] *Email*: gr69@st-andrews.ac.uk. Thanks to colleagues, including Deb Brown, Chris Cordner, John Flett, Kai von Fintel, Rohan French, Anil Gupta, Sally Haslanger, Lloyd Humberstone, Karen Jones, Bernhard Nickel, Graham Priest, Dave Ripley, François Schroeter, Laura Schroeter, Mandy Simons and Shawn Standefer, and audiences at Melbourne, the Graduate Center at CUNY, Pittsburgh, UConn, MIT and UQ for helpful feedback as I worked through this material. This research was supported by the Australian Research Council, through Grant DP180103687.

Generic judgements, like any other declaratives, can be the locus of disagreement. We can disagree about whether birds fly or whether women are emotional. We can use these judgements as a premise or as a conclusion in our reasoning. However, what could count as decisive evidence for or against a generic judgement is by no means clear. Generics display wild and unruly inferential behaviour. For example, most would judge it to be true that *birds lay eggs*, and that *birds that lay eggs are female*, but would balk at accepting the generic *birds are female*, even though there are more female birds than birds that lay eggs. We argue about and reason with generic judgements, but the norms for doing so are opaque to us.

Nonetheless, we *do* argue about generics. If someone makes the claim that men are violent towards women, it is a common refrain to hear the defensive retort "not *all* men," as if that were an objection to the claim. Generics play a significant role, particularly, in the ideologies of our social worlds, of characterising different social kinds and expressing our default orientations toward them (Haslanger 2011; Leslie 2017), and towards ourselves as members of those kinds. So, understanding the affordances of generic judgements, and they way they play a role in our thought and our talk is vital if we are to understand how harmful stereotypes are found, and what we might do to dismantle them.

My contribution in this paper is to show how an inferentialist account of the semantics of generics can explain their pervasiveness, and account for their wild behaviour and their nefarious potential, in a way that is more clarifying and helpful than more traditional truth conditional accounts (Leslie 2012; Nickel 2016; Pelletier and Asher 1997; Sterken 2017; Stovall 2019). The distinctive contribution of this paper is applying this connection to the social phenomenon of the formation of the common ground in discourse. This will help us account for how we come to adopt characterising generics, even when they are not the explicit topic of discussion. We accept generics by a well-understood phenomenon of discourse accommodation (Beaver and Zeevat 2007; Roberts 2015; Simons 2003), applied to *inference*. Understanding the connection, between accommodation, inference and characterising generics will help us understand what is at stake when we argue about generic claims, and will give us better tools to *improve* their use.

## Characterising generics

Generic claims have the form *Fs are G*, where there is no explicit quantifier expression. So, claims like *some men are violent*, and *all men are violent* are not generic claims, and neither are claims like *most men are violent* or *normal men are violent*. In this chapter, I will focus my attention on bare generic claims of the form *Fs are G*, though there is a closely related form with definite or indefinite descriptions *the F is a G* (e.g. *the man is violent*) or *an F is a G* (*a man is violent*).

Our focus is on so-called *characterising* generics, in which items of class *F* are characterised as generally having feature *G*. Each of the examples we have seen so far have this function, but not all claims of the form *Fs are G* characterise individual *F*s in this way. Consider *violent acts are widespread*, and *children ate all the pizza*. These claims do not characterise individual violent acts as widespread, and no individual child is said to have eaten all the pizza. These *class* generics describe the behaviour or features of a group as a whole, rather than their members individually. It will be enough for us to be going on with to focus on characterising generics, because these play a significant role in our thought and talk (especially when it comes to social kinds), and the most difficult issues concerning the truth conditions of generic claims concern characterising generics.

So, to clear the ground, when we consider characterising generics, such as

- *Men are violent*
- *Cows are food*
- *Women are emotional*
- *Mosquitos carry Ross River Fever*
- *Tall people with back injuries find it difficult to drive small cars*

each of which have the from *Fs are G*, and each characterises members of the class *F* as having feature *G*. In each case, to make the generic claim you do *not* thereby say that ***all*** *Fs are G*. (It is true that birds lay eggs, but male birds to not.) You do not merely say that ***some*** *F*s are *G*. (*Some* men are gentle and are not violent. This does not necessarily justify the generic claim *men are gentle*.) The generic claim *Fs are G* does not even require that ***most*** *F*s are *G*, or that ***normal*** *F*s are *G*. (Most mosquitos do not, in fact carry Ross River Fever (RRF). Male mosquitos do not, and neither do mosquitos outside Oceania. These mosquitos are, nonetheless, normal.) Finally, while it might be correct to say that mosquitos are the *kind* of thing that carries RRF, this will only help in specifying the truth conditions of characterising generics when supplemented with an account of kinds and their properties.

As the last example in our list indicates, we can form characterising generic claims with complex predicates. It is one thing to think that semantic and cognitive competence requires some sort commitment to a metaphysics of kinds such as *cows*, *men*, or *tall people*. It is another for that account to extend to complex kind terms such as *tall people with back injuries*, or *Dundee residents who are sympathetic to independence but who weren't included in this survey*.[2]

[2] These examples show that while kind terms and naming conventions facilitate the use of generics to characterize the kinds we name, we can and do use characterising generics not only with kind terms but all

It is important to underline some of the distinctive behaviour of characterising generics. First, they do seem to be truth apt. It is not only meaningful to assert *birds lay eggs*, it is also meaningful to make claims like these–

If *birds lay eggs* then birds and reptiles share at least one feature.

If *that* talk was representative, then *logic talks are boring*.

–so we can meaningfully use characterising generics in the antecedent or in the consequent of conditionals.[3] For any generic *Fs are G* one can form the *opposite* generic, *Fs aren't G*, though it is less clear that we have a natural way to assert the *negation* of a characterising generic, other than the stilted phrasing *it's not the case that Fs are G*. Consider the characterising generic *birds are female*. This seems untrue (to me, at least), when understood as a characterising generic, since birds are sexually dimorphic. The opposite claim *birds are not female* (when also understood as a characterising generic) seems equally untrue, and this seems good evidence that *Fs aren't G* does not express the negation of *Fs are G*. The ungainly wording *it's not the case that Fs are G* (or something like it) seems to be required if we wish to state the negation of the characterising generic.

The 'logic' of characterising generics, such as it is, is complex. We have already seen that *Fs are G* does not mean that *all F*s are *G*, or that *most F*s are *G*, or even that *normal F*s are *G*. As one example of the strange behaviour of characterising generics, even though *birds lay eggs* is true, and *all birds that lay eggs are female* is true, it does not follow that *birds are female*. In other words, from *Fs are G*, and *All Fs that are G are H*, it need not follow that *Fs are H*.[4] In characterising generics, we cannot, in general, *weaken* the consequent.[5] It follows from these and related

---

the different ways we can classify and describe the world around us. Neufeld's essay in this volume for more discussion of this point (Neufeld 202+).

[3] Unlike imperatives, which cannot be used as the antecedent of a conditional, as one example. It makes sense to ask "pass the milk" but no sense to say "if pass the milk then …"

[4] This means that the characterising generic acts significantly different from any *any* account which posits the truth conditions of *Fs are G* in terms of the existence of some selected sub-class of *F*s (whether *most F*s, *normal F*s, *contextually chosen F*s – where the class is chosen independently of the selection of the predicate *G*), such that all *those F*s are *G*, for if *all G*s are *H*, then any such selected collection of *F*s that are each *G*s will also be *H*. Any treatment of *Fs are G* by way of a contextually restricted quantifier will need to allow for the characterising predicate *G* to help determine the selected restriction.

[5] This makes generics even stranger than counterfactual conditionals. If a counterfactual conditional $A > B$ is true (that is, if *A* were the case, then *B* would be the case), and if *B* entails *C* (that is, in *any* possible scenario in which *B* holds, so does *C*), then it follows that $A > C$ is true too (that is, if *A* were the case, then *C* would be the case). The analagous rule fails for characterising generics.

considerations that stating the truth conditions for characterising generics is a very complex matter. (See below for a further discussion of this point.)

The semantics and logic of characterising generics is not straightforward. Let's turn to inferentialism, to see how we might account for this distinctive behaviour.

## Inference

*Inferentialism* is an approach to semantics that attempts to give an account of the semantics of expressions in terms of *norms of inference* (Brandom 1994, 2000). Inferentialism, in this sense, is a species of the wider genus of *Normative Pragmatics*, which encompass approaches to semantics that centre on *norms of use*, perhaps including inference, perhaps not (Lance and O'Leary-Hawthorne 1997; Kukla and Lance 2009; MacFarlane 2010). A normative pragmatic account of the semantics of some concept will be a different way of characterising its meaning, in contrast to a truth-conditional semantics. For a straightforward example, an inferentialist account of the semantics of *conjunction* may characterise conjunction as follows: from the conjunction of *A* with *B* you can infer *A* and you can infer *B*. Conversely, from *A* and *B* together, you can infer the conjunction of *A* and *B*. This is an inferentialist account of the semantics of (sentential) conjunction. A truth-conditional account of the semantics of conjunction goes as follows: the conjunction of *A* and *B* is true if and only if *A* is *true* and *B* is *true*.

It should be clear from this toy example that an inferentialist account of the semantics of conjunction and a truth-conditional account of the semantics need not be in conflict. One might endorse both as correct (Restall 2005, 2009). It is a further question which, if any, of these two kinds of accounts might have priority in any order given order of explanation. However, since the truth conditions of characterising generics are opaque, and it would be surprising if we could explain our competence with generics in terms of our *grasp* of those conditions,[6] and since generics are something we can *communicate* with and *coordinate* on, there is scope for looking elsewhere for an explanatorily rich framework for understanding generics, and what is at stake when we disagree about them.

Since inference plays a central role in any inferentialist semantic explanation, we should spend some time clarifying what inference involves. It will be *very* important to distinguish the material notion of *default* inference from any *formal, deductive logic* understanding of the term.

---

[6] In particular, we should attempt to explain the fact that we acquire generics very early in our cognitive lives, plausibly significantly earlier than we acquire the 'simpler' logically precise universal and existential quantifiers (Leslie 2008, p. 21), which have much simpler truth conditions.

To get an initial sense of the significance of inference, consider the difference between (1) and (2) below:

(1) Tweety is a bird. Tweety flies.

(2) Tweety is a bird. So, Tweety flies.

In both cases, the two claims are made, that Tweety is a bird, and that Tweety flies. But in (2), the second claim is *inferred* from the first. Here, the additional component is that the first claim, to the effect that Tweety is a bird, is presented as *justifying* the second claim. An inference, here, is an action, in which a claim (here, the conclusion, to the effect that Tweety flies) is "backed up" by putative evidence (here, the premise, to the effect that Tweety is a bird). In this example, the premise came before the conclusion, but there is no need for the two components to be presented in that order. We could just as easily have said:

(3) Tweety flies, since Tweety is a bird.

thereby presenting the fact (if indeed, it is a fact) that Tweety is a bird as justification for the claim that Tweety flies. As far as public speech, dialogue, or written text goes, inference can be presented in either direction. There is no requirement that the conclusion *conclude* the inference in any temporal sense.

What goes for public speech can also go for private thought. In my own reflection, I can come to think that Tweety is a bird and conclude *from this* that Tweety flies. Or I could come to believe that Tweety flies, and then wonder if that is really the case, and reassure myself about this fact, since I also know that Tweety is a bird. We can equally understand inference as a kind of transition in our *talk*, and as a transition in our *thought*. In what follows, I will not take a stand as to which of these options, if any, is the fundamental or original notion of inference. Everything in this account of generics will be consistent with the view that inference in thought is *internalised* public inference, and that the norms governing inference are fundamentally a social, communicative matter.[7] And, everything in this account is consistent with the view that we can think of inference as presented in language as *externalised* representation of a more fundamental notion of inference in thought.[8] Or, you could think that this distinction is somehow ill-posed, and that neither a thought-first nor a talk-first account is correct. We need not take a stand, and I draw attention to it merely to clarify that the 'inferentialism' important

[7] This is the normative pragmatic inferentialism of Robert Brandom (2000).

[8] Paul Boghossian's account of inference (Boghossian 2014), as one example, takes it that inferential transitions in thought are fundamental. Similarly, Gilbert Harman's influential account of reasoning in *Change in View* (1984) is all about the dynamics of change in *beliefs*.

to this account is a thin one, according to which inference (whether understood thought-first or talk-first) can play a role in giving an account of the meanings of characterising generics.

To make out this case, it is important to distinguish our target notion of *inference* from *proof, logic* or *deductive reasoning*. The inferences recorded in (1), (2) and (3) above are not *proofs*, and nor need they be treated like proofs by the one who infers in that way. They are not logically valid, and they need not count as deductive reasoning. The inference recorded here is *invalid* in the logical sense, because it is possible that 'Tweety is a bird' be true while 'Tweety flies' is false. Even if we impose some realistic semantic constraint on the interpretation of the predicates 'flies' and 'is a bird' the mere possibility of non-flying birds (emus, hatchlings, injured birds, etc.) ensures that counterexamples to the inference abound. Nonetheless, inferences like these– in which we make everyday transitions from premises to conclusions, despite the possibility of counterexample–are made everyday, and are a part of competently using concepts such as *bird* and *flies*. The relevant notion inference for our account is *contingent* and *material*, and neither *necessary* nor *formal*.

Closely allied to the material nature of inference, is its *defeasibility*. Take the inference from 'Tweety is a bird' to 'Tweety flies'. Suppose that I, not knowing what kind of creature Tweety happens to be, make a justification request for your claim that Tweety flies. You say, 'Tweety is a bird' to discharge that request, thus making the inference. I am well within my rights to take that justification request to be met, and to carry on with the conversation. If you did not finish by saying that Tweety is a bird, but added '… in fact, she is an emu', I would be well within my rights (knowing, as I do, that emus do not fly), to take it that my justification request is now no longer met. What counts as a good inference in some context may be defeated in the presence of new information.

This should not surprise us, if we understand the practice of asking for reasons as imposing some degree of quality control our thought and our talk. There is reason to take up claims from others and from own thoughts and to commit to them for ourselves (it is hard to see how we could get going in our own cognitive lives were we to do otherwise) and there just as much reason to attempt to ask some questions concerning claims that seem, well, *questionable*. Any process like this will look something like the making of justification requests, in the way I have sketched it out here.[9] It is not at all surprising, given the role of quality control, that what counts as meeting a justification request for a questionable claim might vary from context to context, and that inferences may be defeasible in the way that we have seen. After all, sometimes a proffered explanation can go too far and raise more questions than it settles, such as when you go on to explain that Tweety is an emu, or that she is a bird with a wing injury.

[9] For an account of dialogue which takes justification requests seriously, seeing it as a kind of quality control over the commitments made by each interlocutor, see Charles Habmlin's *Fallacies* (1970, chap. 8).

None of this means that the original explanation (that she is a bird) would not have done enough to assuage us. In wanting our questions answered, we do not necessarily need *every* question to be answered. We can treat a claim as justified without meeting the logicians' standard of exceptionless deductive validity.

Nonetheless, we can see why the notion of inference can lead to the logicians' sense of validity as a kind of limiting case (Girle 2016; Dutilh Novaes 2015, 2016, 2020), since there are two different ways that the answer to a justification request can fail. Your answer to my request for a justification can be rejected as a claim that is itself *untrue*, or itself *unjustified* (say, I ask you to justify your claim that Tweety is a bird). However, I can grant that your further claim is in order, but reject it as not meeting the claim to justify the conclusion. I could be the skeptic and say, I grant that Tweety is a bird, but isn't it possible that nonetheless, Tweety might not fly? The logicians' sense of *deductive* validity is an attempt to mark out a kind of limiting case of good inference, where the offered justification leaves *nothing* out, and anyone who accepts the premises but still takes the conclusion to be in question has shown that there is some kind of failure to communicate, rather than a disagreement whether to say 'yes' or 'no' to a shared question.[10] For our purposes, however, we need not concern ourselves with the idiosyncratic case of deductive logic.

Before turning to the connection between generics and inference, let's pause to notice the breadth of uses we find for our inferential capacities. The 'quality control' nature of making justification requests can be applied, usefully, in other domains. An important one for creatures like us is in our *planning* and *decision-making*. In planning for the future, we consider different 'possibilities'. Considering a possibility need not be understood as a kind of observation, where we come across a thing called a possibility and attempt to describe it as best we can. Instead, we can understand planning as involving inference from suppositions. When we consider what we take to be a future possibility, we *suppose* it comes about, and infer from there, as best we can. We apply our practice of inference, drawing out consequences, not to something we take to be the case, and not to something someone else has asserted. Instead, we apply the same reasoning processes, asking ourselves what else would follow *were* that possibility to come about. In this way we can do some thinking in advance, and make decisions now, to apply in future circumstances, if they arise, or use our understanding of the significance of future possibilities to inform our decisions about options to take (Lance and White 2007). Our inferential capacities are needed in each of these ways of engaging in with the world, and so, given the connection between characterising generics and inference we will find below, it will turn out that generics, too, play a fundamental role in our planning and deciding.

---

[10] If I come across someone who grants *A* together with *if A then B* but still thinks that *B* is in question, then it seems that they must use "*if … then …*" in some way that I do not yet understand (Carroll 1895).

An inferentialist account of the significance of characterising generics does not start by giving the truth conditions of a characterising generic statement. Instead, we explain its semantics in terms of *inference*. This seems like a natural thought, for it is widely recognised that there is some kind of connection between generics and default, defeasible inference (see Leslie 2012; Pelletier and Asher 1997 for example) and default, defeasible inference plays a central role in inferentialist semantics (Brandom 1994, 2000). At the very least, it is natural to think that while the inference from *Fs are G* and $Fa$ to $Ga$ is not deductively *valid*, the inference is appropriate as a *default*. This seems to be the natural way to characterise the inferential power of making the claim *Fs are G*. So, one governing rule for a characterising generic is this: from the assumption that *Fs are G* and the assumption $Fa$ we can defeasibly infer $Ga$. This seems appropriate at first glance: given an arbitrarily chosen bird, it is appropriate to defeasibly infer that it flies, and if I take it that men are violent, it might seem appropriate to infer, at least defeasibly, of an arbitrarily chosen man, that he is violent.[11]

When it it appropriate, though, to make such a generic claim? The natural inferentialist answer is the converse: when we can make the default inference from $Fa$ to $Ga$, when you make no special assumptions about $a$. This is the simple and direct inferentialist account of the meaning of characterising generics, and it has the same form as other kinds of inferentialist semantics for concepts such as conjunction, the conditional, or explicitly marked quantifiers.[12] Though this account is relatively natural, it has faced some tough criticism, especially from those, like Pelletier and Asher, who take the connection between generics and default inference seriously.

> … default logic does not provide us with an acceptable formalization of generic statements. Default rules are rules, and therefore are sound or unsound – rather than sentences, which are either true or false. If we analyze characterizing sentences using default rules, these sentences would not have truth values, and their meanings could not be specified by an ordinary semantic interpretation function. One consequence of

[11] As we will see below when we discuss low-rate generics such as *mosquitos carry RRF*, it will be important to keep in mind that inference goes in two directions. We can infer some new conclusion *C* from a background *B*, or we may discharge a request to justify or explain *C* from that background context. Characterising generics are connected to inference in both these directions.

[12] In particular, it has the same form as the inferentialist account of the universal quantifier. What can you conclude from $\forall x\phi(x)$? You can prove $\phi(t)$, for any singular term $t$. When are you in a position to prove $\forall x\phi(x)$? When you can prove $\phi(a)$ for some *arbitrary* term $a$. This is not spelled out by way of truth conditions (in particular, $\forall x\phi(x)$ might turn out to be true even when we cannot *prove* $\phi(a)$ for an arbitrary term $a$. Nonetheless, the soundness and completeness results proving the equivalence of validity as defined by the inference rules and model-theoretic semantics shows that there need be no conflict between truth conditions and inferentially characterised conditions, even when they take radically different form (see, e.g., Restall 2022).

> being neither true nor false not being in the language is that characterizing sentences would therefore not talk about the world, instead they would talk about which inferences to draw. And this seems to us to be a strike against such an account. (Pelletier and Asher 1997, 1152)

It should be clear that this criticism is misplaced if used to target *this* inferentialist account of characterising generics. The idea is not that a generic claim *Fs are G* is *true* if and only if we can make the default inference from $Fa$ to $Ga$. That would be to use the inference rule as a *truth condition*, but the inference rule does not specify truth conditions. The inference rules tell us when we can infer something *from* a characterising generic, and when we can *conclude* a characterising generic. Like the other sentences of the language, characterising generics are truth apt simply because they are, on this account, the kind of things that can feature as premises and conclusions in reasoning. The inference rules do not take the form of truth conditions.[13] It is not my place to give an extended defence of the inferentialist semantics for characterising generics. Thankfully, that work has been done by Preston Stovall in his "Characterizing Generics are Material Inference Tickets: a proof-theoretic analysis" (Stovall 2019). In the remaining sections of this chapter, I will take the broad brushtrokes of this view as given, and show how the connection between default inference, the common ground and accommodation phenomena can help explain the difficulties that arise around disagreement involving characterising generics, and how we might, nonetheless, make progress.

* * *

So, with this understanding of inference and its role in discourse and thought in view, we can more clearly understand the connection between default material inference and characterising generics. The key thought is that default inference is acquisitionally and conceptually prior to our characterising generics. The force of a characterising generic *Fs are G* is to license the default inference from $Fa$ to $Ga$. If I grant *birds fly*, I ought not demur at the default inference from *Tweety is a bird* to *Tweety flies*. Conversely, if I accept the default inference from $Fa$ to $Ga$ (where I appeal to no particular features of $a$ other than the assumption that $Fa$), then I am in a position to grant the characterising generic *Fs are G*. This does not mean that the characterising generic is *about* whatever inferences we accept, since these statements about the norms for

[13] The distinction is most stark in the case of the natural deduction inference rules for the logical quantifiers. These rules *clearly* characterise the semantics of the quantifiers ∀ and ∃. They do not do so by taking the form of truth conditions, because they do not involve any characterisation of a domain of quantification, or make use of assignments of values to variables, or any of the parephenalia of models for first-order predicate logic which are the appropriate way to characterise the quantifiers truth-conditionally. Despite that, the inference rules do manage to inferentially characterise the behaviour of the quantifiers in a way that agrees with the truth-conditional account. The two accounts differ, but are not in any conflict.

granting generics, how we are in a position to infer to a generic, and what we can infer from them, and are not statements of their truth conditions.[14] It should not come as a surprise that an account in terms of inferential power should itself involve the inferences that we license.

The power of this analysis is that it explains the key distinctive behaviour of characterising generics. Characterising generics have exceptions (birds fly, but penguins do not, despite being birds), in just the same way that default inference has exceptions. We have an explanation of *why* default inference comes with exceptions: it would be a waste of time and energy to demand that all justification requests be met only at the highest, 100% standard. We count justification requests as met with a much lower threshold, accepting reasons that we take to count in *this* case, even though they might not count in *every* case. Since granting an inference is taking a justification request (whether implicit or explict) to be met, then we have an explanation of why it is that default inference carries exceptions. Since generics inherit features from default inference, we can see why it is to be expected that characterising generics have exceptions, too.

Let's consider another example of strange behaviour of generics: the failure of weakening of the consequent, mentioned earlier, when we introduced generics. We grant *birds lay eggs* and we grant, also, that *all egg-laying birds are female*. It does not follow, though, that *birds are female*. We resist granting this characterising generic, despite acknowledging that all egg-laying birds are female. Why do we so resist? This can be explained by appealing to how we treat justification requests and inferences. The inference from *Tweety is a bird* to *Tweety lays eggs* is acceptable, at least in some wide range of contexts. If I am unsure of what kind of animal Tweety is, and I am surprised that she has laid an egg, you can explain that she's a bird, and then my surprise is addressed. If the relevant options are that Tweety is a bird, or a mammal, or a reptile or some other creature, then settling on *bird* helps settle the issue between egg-laying and live birth. It is *much* harder to find a context in which making the inference from *Tweety is a bird* to *Tweety is female* is anywhere near as helpful. In most contexts in our thought or our conversation, we are not going to meet a justification request for the claim that Tweety is female by granting or learning that she is a bird.[15]

---

[14] See the earlier discussion of the natural deduction rules for the logical quantifiers. These rules state that to *prove* $\forall x \phi(x)$ from some set of assumptions, I can prove $\phi(a)$ for an *arbitrary* $a$ (a name '$a$' that does not appear in the assumptions used). This rule characterises the universal quantifier, and gives it a meaning that gives it the usual truth conditions. So, it need not follow that these inference rules give the universal quantifier deviant meaning that is somehow 'about' arbitrary names, or provability from premises.

[15] This is not to say that you cannot cook up artificial contexts in which the inference from *x is a bird* to *x is female* makes sense. Take a background in which we are considering the animals in some large factory farming facility. There are humans, birds, and maybe some domestic dogs. The birds are all battery farmed chickens, and these are all (or are mostly) female, caged for their eggs. In this context, where I consider a randomly chosen creature in this facility, it makes sense for me to make the inference from *x is a bird* to *x is*

Let's consider the phenomenon of low-rate generics with striking properties. Even though very few mosquitos carry Ross River Fever (RRF), it makes sense for us to grant *mosquitos carry RRF*. One explanation may be that these are not, at root, characterising generics in which we describe individual mosquitos, but a *class* generic, where we ascribe a property to the whole population of mosquitos. This population carries RRF, in the same way that this population is *widespread*. However, the cases are also dissimilar, in that individual mosquitos *do* carry RRF, so let's see whether an inferential treatment of these cases can make sense of this phenomenon, if we understand the claim as a characterising generic.

While it might be odd to make the inference from a given claim *x is a mosquito* to *x carries RRF* in any given context in which an arbitrary mosquito is under consideration and we are thinking of the inference as progressing forward from premise to conclusion, there are important contexts in which this inference *does* make sense, specifically when we reason backwards, and ask questions concerning how it is that someone has caught RRF. We can answer a justification request for "I caught RRF" by "I was bitten by mosquitos", and this is not only an acceptable answer in many conversational contexts, it is also one that carries salient and actionable information. Similarly, if we are concerned about RRF, and wish to know how it is transmitted, it makes sense to be told that it is *mosquitos* that carry RRF, rather than flies or midges. This explanatory and justificatory connection makes sense, even when the rate at which the property is held in the population is quite low, especially in cases like this where we are keen to protect ourselves from danger.[16] Given the importance of making the inference in cases such as these, we can see why it can make sense to grant the generic *mosquitos carry RRF*.

Another striking feature of generics is the difficulty we have in expressing a denial of a generic claim. Consider, for contrast, the universal quantifier. To deny the claim *all Fs are G* it suffices to assert *not all Fs are G*. Consider a comparable assertion that expresses a denial of the generic *Fs are G*. This is not done by way of the opposite generic *Fs aren't G*, at least if that has the meaning that, generically, *F*s fail to have the property *G*. After all we can reject both *Fs are G* and *Fs are non-G* when rejecting both inferences, from something being an *F* to its being a *G* and to its not being a *G*. One can reject the generic *birds are female* and *birds aren't female* equally. How can we express this rejection by way of asserting a negative statement? The clearest way to do so is something like *it is not the case that birds are female*, where the negation clearly takes wide

---

*female*. In just the same way, it is appropriate *in these contexts* to grant *birds are female*, where this is understood as appropriately contextually restricted, to when we are talking only about the birds *around here*. The fact that generics seem *as* appropriate as the corresponding inferences in contexts such as these seems to be evidence in favour of the generic/inference connection.

[16] See Sarah-Jane Leslie's "The Original Sin of Cognition" (Leslie 2017) for more on the phenomenon of low-rate striking property generics.

scope over the generic statement, but even this can be plausibly understood as an inner negation, claiming that *birds aren't female*. Given the difficulty of fixing scope with generic expressions, it can be much clearer to make some quantifier expression explicit, to say *many birds aren't female* to indicate why we resist the generic *birds are female*, but as we have seen in the case of low-rate striking generics such as *mosquitos carry* RRF, or in the response *not all men* to accusations to the effect that men are rapists, holding to generic can resist the presence of an explicitly quantified negative claim in response, so this strategy cannot be relied upon to provide a clear statement of an opposing position. This striking behaviour of negated generics makes disputes in which generics are the subject matter particularly intractable.

Since the "logic" of generics seems wild, any reasonable account of their truth conditions turns out to be complex. This is a matter of immense theoretical interest, on which there is very little agreement (Leslie 2008; Liebesman and Sterken 2021; Nickel 2016, 2017; Pelletier and Asher 1997; Sterken 2017). Here is one example from the recent literature, in a paper by Sarah-Jane Leslie (Leslie 2008, 43):

> Though there may be a further refinement or two needed, we can describe the circumstances under which a generic of the form '*K*s are *F*' is true as follows:
>
> The counterinstances are negative, and:
> If *F* lies along a characteristic dimension for the *K*s, then some *K*s are *F*, unless *K* is an artifact or social kind, in which case *F* is the function or purpose of the kind *K*;
> If *F* is striking, then some *K*s are *F* and the others are disposed to be *F*;
> Otherwise, almost all *K*s are *F*.
>
> I would suggest that these worldly truth specifications—these descriptions of how the world must be for the sentence to be true—should not be mistaken for semantically derived truth conditions, however.

The complexity of these truth conditions is due to the wild behaviour of generic claims. The reasons for each component in the truth conditions need not detain us here. However, we should note that characterising generics with complex kind terms like *tall people with back injuries* and *Dundee residents who are sympathetic to independence but who weren't included in this survey*, put pressure on the division into artifactual, social and natural kinds, and as Leslie hints, her account may need to be refined to deal well with generic expressions with complex kind terms. It follows from this that the connection between the use of generics and the propensity to *essentialise* categories is quite subtle, and worth further reflection (Haslanger 2007, 2011; Leslie 2017).[17]

[17] I note here that the analysis of generics given here is completely orthogonal to the question of whether generics in any way *essentialise*. The question will become: whether and when–and how–do inference and

As Leslie points out, these are "worldly truth specifications"–accounting for what the world must be like to satisfy the characterising generic *K*s are *F*–and are not to be taken to be *semantically derived* truth conditions, which in some sense mirror the structure of the ground-level claim. Leslie uses the following example to illustrate the point. The claim that *Bob is red* may be true if and only if Bob is experienced as red when observed by standard observers in standard conditions. This is a statement of worldly truth conditions for the claim. However, the *semantically* derived truth conditions will not have this form, because semantic competence with the claim *Bob is red* need not involve any commitment concerning standard observers or standard conditions. The semantically derived truth conditions may have the form of an object (in this case, the referent of *Bob*) falling under the extension of a predicate (in this case, the extension of *red*). Leslie argues that the *semantically* derived truth conditions for characterising generic claims may simply *use* the generic quantifier. When it is true that *Fs are G*? If and only if, $(Gen\ x)(Fx, Gx)$, where '$Gen$' is the generic quantifier, used in the metalanguage.

In one sense, such *deflationary* truth conditions for generics must be correct (at least, they must be, if characterising generics have truth conditions at all), but they do not help when it comes to explaining competence with generic expressions, or what capacities one has to employ in order to be able to use characterising generics. If the truth conditions for generics are either too complicated to be able to account for competence (as Leslie's example worldly truth conditions are), or too *thin* to be informative (as the deflationary truth conditions are), and it is for this reason that the inferentialist account has explanatory power.

In the following section, we will see how the explanatory power of the inferentialist account extends further, when we see that we do not just *learn* generics when we come to consider them for ourselves and decide whether we agree with them or not. Instead, we can find ourselves committed to generic claims by a kind of conversational osmosis from our surrounding community. It is to this phenomenon that we now turn.

## Accommodation

Conversation is a shared exercise, in which the participants build something together. A part of what they build is the *common ground* of that conversation, those things that everyone in the conversation has granted, and from then on, can take for granted. The notion of common ground has been central in the contemporary study of the interaction between semantics and

---

explanation essentialise? One way to address this question is to consider the range of different *contexts* in which an "inference ticket" is to be used. The whole point of learning a generic is to remind us that an inference is available, and available in a range of future contexts. Just how far can and should that range go?

pragmatics in recent decades (Geurts 2018; Haslanger 2011; Simons 2003; Stalnaker 2002). Kai von Fintel describes the common ground as follows:

> The COMMON GROUND of a conversation at any given time is the set of propositions that the participants in that conversation at that time mutually assume to be taken for granted and not subject to (further) discussion …
>
> When uttered assertively, sentences are meant to update the common ground. If a sentence is accepted by the participants, the proposition it expresses is added to the common ground. (Fintel 2008)

What is shared between the participants of a discourse includes more than a body of propositions. Since conversations involve questions, we also keep track of a shared list of questions currently under discussion (Beaver et al. 2017; Roberts 2012, 2018), and if the conversation includes imperatives, we will also together keep track to-do list for each individual (Portner 2004, 2009). Furthermore, the fact that conversation is a shared enterprise does not mean that each participant's own private commitments are transparent to the others. There is a difference between the public common ground (what we have committed to in public, by way of what we have said and what we have let stand) and our own private commitments (Murray and Starr 2018). A public record of a common ground is not inconsistent with deceit (where we assert something in conflict with what we hold true, for the purposes of manipulating others) and neither is it inconsistent with silent dissent, in which we let what someone else says stand, without raising an objection, even though we privately disagree. In both cases, what is *said* in the conversation can stands as the common ground of that discussion, despite this not being a record of the private commitments of the participants.

One important feature of the common ground, in understanding the dynamics of discourse, is the phenomenon of *accommodation*. The shared propositional state of the conversation – those claims that we together take for granted – can shift not just by updating with the content of what any participant explicitly says. Other items can enter the common ground in order to accommodate what has been said. If, after giving a presentation as a visiting scholar at a research seminar, I say *if there is dinner after the seminar, my son will come*, then, if that statement meets with no objection, the common ground of that conversation will be updated, not only with the conditional claim, but also with the information that I have a son, and most likely, that I have *one* son. I did not literally say that I have a son, let alone that I have only one. Speaking literally (if unhelpfully), the content of what I said could be true even if I had been childless, had it been certain that there was no dinner after the seminar, and so the issue of anyone coming with me would not arise. Of course, no-one would interpret such a statement in that way, and we all take it for granted that I have the son who I mentioned in the claim. If my claim passes into the common ground with no objection, then another participant in the conversation can meaningfully ask *how old is he?*, and it is obvious to all concerned that the question is about my

son, previously mentioned. This phenomenon is called *presupposition accommodation*. The details of *how*–and *which*–presuppositions are accommodated is a matter of active research (Beaver and Geurts 2012; García-Carpintero 2016; Simons 2006; Roberts 2015), and we need not go into the details of any account of presupposition accommodation here. Suffice it to say that it is now very much a live issue to explore the dynamics of conversation and the way that our commitments are updated not only by way of the content of what is explicitly asserted but any number of ways that are less explicit. This area of research has also included rich connections to matters of ethics and normativity more broadly construed (Langton 2015; Adams 2020).

Given the connection between characterising generics and inference delineated above, we can see that generics are also added to the common ground in conversation, in ways that are reminiscent of accommodation phenomena, but have their own distinctive features. The point is a simple one. On the standard picture, if I say *Tweety is a bird. So, she flies,* in conversation, and I meet with no objection, then the common ground is updated with the information to the effect that Tweety is a bird, and that she flies. However, my statement did more than just make the assertion that Tweety is a bird, and the assertion that she flies. In addition, I treated the connection between these claims as an inference. The inference *itself* could meet with resistance, and be rejected in its own rights, without rejecting either assertion. It makes sense, then, to think of the inference as being recorded, somehow, in the common ground, to mark the difference in our conversational commitments in the case where an inference is rejected from a case in which it is not. We have a candidate for how the common ground can be updated in cases where we make the inference from $Fa$ to $Ga$ and meet with no resistance: it is that the common ground is updated with the generic *Fs are G*. At least in this local conversational context, if we grant the inference *from* $Fa$ to $Ga$ (so, in this conversation we appeal only to $Fa$ to answer the justification request for claim $Ga$), then we have granted, implicitly, the generic *Fs are G*. If, in a conversational context, we can appeal to $Fa$ alone as a justification for $Ga$, then at least relative to this context, we are committed to the corresponding generic.

If this view is correct, then the common ground updates with the generic *Fs are G* whenever an inference from $Fa$ to $Ga$ is admitted without objection. This is like presupposition accommodation in the way that a propositional content enters the common ground without being directly asserted. However, the propositional content that is admitted (the characterising generic) is not some presupposition of something else asserted. It is the propositional content that corresponds to the making of inference itself. The inference from $Fa$ to $Ga$ is licit, in this context, if (in this context) *Fs are G*. The proposition to update is directly represented in the conversation by the inference (the '*so*', '*therefore*' or however else the inference was represented), and there is no complicated negotiation to expand the common ground appropriately to meet the presuppositions of some other asserted content. The common ground updates in a manner that is explicitly represented in the discourse itself.

However, although the points at which inferences are made in discourse are explicit, and can be questioned and rejected just as assertions are refused admission into the common ground, blocking an *inference*, in particular, is not straightforward. If I say $Fa$, *so* $Ga$, it is easy enough for you to block either of the assertions of $Fa$ and $Ga$ on their own: you simply call either assertion into question with an 'are you sure?' or 'I don't agree'. It takes more finesse to single out the *inference* as faulty, especially when you grant $Fa$ and $Ga$ on their own terms, but you take it that $Fa$ doesn't count in favour of $Ga$. Take an example circumstance where someone from a minority (say, a logician, in the Philosophy Department) has given a boring talk, and a conversation partner says 'of course it was boring, it was a *logic* talk', then in the context where the talk did happen to be boring, and it is not in question that it was, indeed, a logic talk, to object to the inference you have to make the connection more explicit than has been said in the inferential transition from one claim to the other, by saying something like 'the fact that it was boring has nothing to do it being a *logic* talk.' Simply saying 'no' to the inference does not have the required precision, because the refusal could equally be interpreted as a refusal to accept either of the assertions instead of as an objection to the inference itself.

Inferential transitions go by quickly in our conversations and in our thought, and as a result, we find ourselves—at least in those contexts where they go by—committed to them, just as we find ourselves committed to the assertions that are made in those conversations. We sometimes *learn* things in our conversations with others. Some of those things we hear and find ourselves agreeing with become standing beliefs as we hold onto those commitments beyond the scope of that interaction. This occurs not only though a process of considered rational reflection on each individual item that we have accepted. Sometimes we simply find ourselves continuing to believe what we were told. It is not surprising that we do so, given the function of conversation and communication, and the way we model our behaviour on others. It is not at all surprising that when we find ourselves in a community who make certain kinds of inferences, we find ourselves inferring in that way, too. Inference and explanation are, after all, partly social practices. As a result, we should expect to find ourselves committed to characterising generics, whether we have rationally reflected on those commitments or not. And upon reflection, we find that this is indeed how we behave, as thinking creatures who orient our selves to the world with views filled with generic connections and the default inferences they express. We come to believe characterising generics by accommodating inferences.

## What we can *do*

Many characterising generics in our conceptual schemes enforce injustice and oppression, as do the inferences they make explicit. Our views of those around us are filled with stereotypes, many of them bearing the marks of our histories as oppressor and oppressed. Our languages and conceptual schemes bear witness to our conflicted histories, and the language of the

dominant group and the way we wield that language and our explanations of the world around us and our place in it cannot help reflect that terrain of privilege and exclusion. This goes down to the level of the individual inferences we accept, as these are the fulcra leading from one concept to another, the tiniest links in our web of connections that represent our expectations and characterisations of what is normal and what is out of the ordinary, of what stands in need of justification, and what can count as a sufficient answer to a given question. The links in this network of interconnections are represented explicitly by the characterising generics we accept. It is not surprising that these claims are become a site of conflict, when we disagree about whether *men are violent, women are submissive,* or *Muslims are terrorists*. There are contested understandings of how the world *is* and how it *should be,* and it is no surprise that when we move far beyond the sharply defined quantifiable expressions that can, on occasion, be decisively settled one way or another (as much as anything is settled in the context of heated debate), to the defeasible wilds of default generalisations that are notoriously resistant to counterexample, this disagreement is well-nigh intractable.

So, what hope is there for critique and reform for contexts in which toxic and oppressive views, expressed in generic terms, have taken root? What tools do we have at our disposal that might help elicit change? In this concluding section, I will make some suggestions and recommendations that come to light when we recognise the connections between characterising generics, default inference and accommodation.

***Denying a generic is not enough***: As we have already seen, arguing over a characterising generic *Fs are G* is true or not is unlikely to shift views concerning them. If the generic is deeply held, then no number of counterinstances, even up to granting that *most Fs aren't G* would be enough to shift commitment to *Fs are G*, as the Ross River Fever cases show us. The reply *not all men* to the generic claim *men are violent* is beside the point, at least when it is thought to serve as any kind of evidence against the original generic claim. This facile reply is so prevalent—and so transparently an attempt to avoid the issue in question—that it became recognised in the form of the #notallmen meme.[18] Replying to a discussion of the normalisation of gender-based violence by denying a *universal* claim (which is all that a #notallX reply ever does, semantically speaking) is not addressing the point in question, and the #notallmen memes mocking those who attempt to avoid self-reflection and avert criticism hit their target, since these rhetorical moves are manifestly self-serving.

[18] See Kelsey McKinney's 2014 *Vox* article for a discussion https://www.vox.com/2014/5/15/5720332/heres-why-women-have-turned-the-not-all-men-objection-into-a-meme.

A natural second response to such unhelpful discussion is to try to avert such responses, by *avoiding* generic vocabulary in favour of more explicit and semantically transparent claims. This strategy has some power.e

***In limited domains it may help to avoid generics***: Instead of making the generic claim *Fs are G*, we can be more explicit, to say *most Fs are G*, or *many Fs are G*—or even better, even more sharply defined claims such as *at least 20% of the Fs are Gs*, or *more Fs are G than non-Fs are*, and so on. A #notallFs response is even more clearly beside the point in cases like these, because the claims explicitly represent the *extent* of commitment made, and this commitment (to a quantity, or relative comparison, etc.) is present as a possible focus of attention. Formulating claims in this way, we move into the mode of explicit quantification and engage in statistical inquiry, with all the rigour that entails. In certain explanatory practices, this makes a great deal of sense. It is clearer to everyone concerned (to a significant degree, at least) what counts as evidence for or against an explicitly quantified claim, and if we have the syntactic discipline to restrict our vocabulary in this way, we can change the arena of dispute and engagement to this new ground.

This strategy has much to recommend itself, and whenever claims can be reformulated without loss to the content under discussion, it is worth sharpening up our claims, whether the generic in question (such as *men are violent*) is one that is useful and we wish to retain and to better understand, or whether it is (like *Muslims are terrorists*) a generic that we wish to dislodge and undercut. In either case, shifting discussion to more explicitly quantificational vocabulary is one way to move to ground at which what counts as evidence *for* or *against* the claim in question is better understood, and fewer unhelpful #notallFs responses are likely to be elicited.

So, this strategy is useful. However, given the analysis of the semantics of generics discussed above, we can see that this strategy cannot provide the *whole* answer to dealing with slippery generics. As we have seen, generics are acquisitionally prior to more explicitly quantified vocabulary. It takes *discipline* to formulate claims with explicit quantification, and we (and our interlocutors) do not always put in the work required to avoid generics. If you (*per impossibile*) belong to an scientific community that speaks only the purest first-order predicate calculus, you still converse with those who lack such discipline, and you need to respond to their claims.

More importantly, though, we cannot avoid the issue raised by the imprecise and messy nature of generics because they are the explicit tip of the implicit iceberg of everyday imprecise and messy default inference, and *that* practice is never going away, even if we decide to not make it explicit in our language with generic vocabulary. We all make default associations and inferential transitions in our thought and our talk, and the clusters of conceptual associations and default connections between concepts lurks in the background, even if we do not make those connections explicit in generic claims.

***The practice of default inference will not go away***: Everyone engages in "reasoning" in the broadest sense of making inferential and explanatory transitions, in our thought and in our talk. This is *not* (and can never be) restricted to the canons of deductive validity—we are committed to default inference. We all make the kinds of defeasible inference steps that ground characterising generics. If I make the default inference from $Fa$ to $Ga$, then even if I refrain from articulating this with the generic expression *Fs are G*, the underlying issue—that I conceive of Muslims as terrorists, or women as submissive, or men as violent, etc.—remains. Even if I have somehow found a way to keep these transitions out of my practice of assertion by being careful to avoid generic expressions, I still licence those inferences, and those connections are salient to me, despite having eliminated generics from my thought and talk. The inferences I grant will remain unperturbed, and no quietism concerning generics will do anything to shift those transitions. Furthermore, whenever I communicate with those around me, I will accommodate the inferential transitions made by others, whether knowingly or not. The default connections between concepts and the fields of semantic associations are communicated, whether implicitly or explicitly.

It follows that doing away with generics and shifting attention only to explicitly quantified claims is only one part of a response to the use of oppressive characterising generics. We must attend to the broader practices of inference and explanation out of which those generics arise. Why might an inferential connection between two concepts arise in the first place?

***We seek connections between salient concepts***: The practice of inferring $Ga$ from $Fa$ will arise in a community only when the concepts *F* and *G* are salient. So, one blunt way undercut the default inference from $Fa$ to $Ga$ is to shift to a practice in which the question of whether something is *G* or not *doesn't arise*. If I do not care about whether something is *G* or not, and do not need to have any understanding of how *G* is connected to other concepts, this would allow for the inference from $Fa$ to $Ga$ to dissipate. If you are in a community in which Ross River Fever is not a problem, you do not form the view that mosquitos carry it. It is only in those communities where people *care* about a concept that the generics become deeply rooted because we wish to understand how to avoid those things we wish to avoid, and attract those things we desire, and in any range of ways, influence our environment in ways that further our goals.

In the case of unjust stereotype generics such as *Muslims are terrorists*, this becomes salient only in communities in which terrorism is a live concern, whether reasonable or not. As a result, it is unsurprising that explanatory and inferential connections take root for concepts such as these.

When we *care* about the property *G*, we won't want to revise the concept away. Another approach to undercutting the *F* to *G* inference is to find an alternate explanation for $Ga$, other than $Fa$. To start with our running example, if not all species or genus of mosquito carry RRF, then we could *refine* our generic to the more specific one, that *Mosquitos <u>of genus Culex</u> carry* RRF,

while being able to also say that *mosquitos outside that genus don't carry RRF*.[19] For communities that are able to keep track of mosquitos by specific genus, a more specific generic like this can easily take root and replace the more general one that characterised all mosquitos indiscriminately. Unfortunately, for a community that has no other use for distinguishing mosquitos by genus, and no ability to keep track of which kind of mosquitos are where, such a conceptual revision is unlikely to take root in an everyday inferential practice, since the finer-grained classification has no grip. In cases like these, it is understandable that the coarse-grained classification and its corresponding characterising generic remains in our conceptual scheme.

To attempt to use this kind of strategy to undercut a politically charged and oppressive generic such as *Muslims are terrorists*, one might attempt replacing them with more useful explanations, that combine (1) a higher degree of explanatory power and (2) which use concepts that are also in wide use in the target communities. One possibility would be to challenge the stereotypical views of terrorists by appealing to research showing that having a history of perpetrating domestic violence is a much stronger predictor than any other factor of involvement in terrorism and mass killing (Pain 2014). The concept of domestic violence is now in widespread community use,[20] and violence in one domain is salient to violence in another, so it is natural to see a connection with involvement in mass killing. Given that a community concerned about mass violence is looking for ways to understand it, alternative explanations using available concepts at least have the *potential* to shift explanations from unhelpful, unjust characterisations in more just and productive directions.

To shift explanatory practices across a community requires *work*. Any shift of this form is not merely replacing assent to one generic with assent to another, but changing how people explain and infer. For a new practice to take root and to be established, participants need to see things in a new way, to make new connections and to avoid old ones. This is an essentially *social* practice, because what counts as a successful explanation or inference in dialogue is one that is accommodated when used, and it is also essentially *normative*, in the sense that default inferences form part of the fabric of our norms for communication. To change social norms is a collective enterprise (Bicchieri 2017). If there is enough resistance, or if there is not uptake in a wide-enough community, an otherwise worthwhile practice may not take root. This is one

[19] Compare Neufeld's advice in this volume (202+) to move away from appealing directly to category terms to more fine-grained explanatory classifications.

[20] Jeff Hearn's *The Violences of Men* (1998) gives an insightful perspective on the introduction of this concept in the United Kingdom through the 20th Century.

reason why these kinds of conceptual shift first occur and establish themselves in smaller subcommunities before spreading to a wider polity.

***New connections can take root in alternative communities***: As we saw above, it took time to introduce the concept of domestic violence into wide circulation. Given that conceptual and communicative practices need time to take root, it is unsurprising that concepts with wider uptake arise from the practice in smaller communities, with their own distinctive ways of looking at the world, of drawing connections, and explaining things to themselves and to each other. Changing norms requires uptake in a network, and the larger the network, the more effort is needed for any change to be established.

Patricia Hill Collins (2000; see especially Chapter 5) gives an account of the different kinds of safe spaces in which Black women were able to come to their own voice. To have a voice, and to be able to use it, is not only to recognise that you have something to say, but to have *uptake*—to recognise yourself as being heard. For the kinds of conceptual connections we have been discussing, this form of uptake is vital. For something to *count* as a reason, it must be *accepted* as one. If I offer a new way of seeing things, it will not spread if others do not follow along. As José Medina puts it (2013; p. 225), these social acts must be *echoable*. They must be repeated, and imitated, if they are to spread, and if the new vocabulary, and new conceptual connections, are to shift a wider practice. Such uptake is more likely to occur at the beginning if the innovators of the new practice do not have to spend all their time clearing the ground and dealing with objections, whether those objections are in good faith, or not.

Members of marginalised groups especially labour under the burden of having to undercut stereotypes and controlling images which maintain the status quo against them (Polhaus 2017, p. 21). There is no neutral public square of discussion, no "objective" discursive space in which reasons can be disinterestedly tested. There is a time and space for descriptive and explanatory practices to be tested with a sympathetic and critical audience—a safe space in which views are developed in smaller community—before venturing out into the wider world.

Given that one way that we *acquire* our explanatory and inferential capacities is in dialogue with each other, if new connections are to emerge onto the public stage, they will take root in smaller subcommunities which develop not only new concepts and new connections between concepts, but also new kinds of salience structures, in which the inferences that are granted have shifted, and new, productive explanations may take root.

The distinctive behaviour of characterising generic judgements—their prevalence in our thought and talk, their robustness in the face of counter-evidence, combined with the difficulty of getting new generics accepted, and their use for unjust social ends—is explained by their grounding in our inferential and explanatory practice. When we pay attention to that practice, we can see that any shift in these stereotypical generic judgements cannot stay at the level of

arguing for or against this generic claim, but must shift the underlying social practice of classification, explanation and reason-giving. It is here that we find possibilities for reforming and revising those practices and the views they represent.

## References


Adams, N. P. 2020. "Authority, Illocutionary Accommodation, and Social Accommodation." *Australasian Journal of Philosophy* 98 (3): 560–73. https://doi.org/10.1080/00048402.2019.1621913.

Beaver, David, and Bart Geurts. 2012. "Presupposition." In *The Stanford Encyclopedia of Philosophy*, edited by Edward N. Zalta. CSLI.

Beaver, David, Craige Roberts, Mandy Simons, and Judith Tonhauser. 2017. "Questions Under Discussion: Where Information Structure Meets Projective Content." *Annual Review of Linguistics* 3 (1): 265–84.

Beaver, David, and Henk Zeevat. 2007. "Accommodation." In *The Oxford Handbook of Linguistic Interfaces*, edited by Gillian Ramchand and Charles Reiss. Oxford University Press.

Bicchieri, Cristina. 2017. *Norms in the Wild: How to Diagnose, Measure and Change Social Norms*. Oxford University Press.

Boghossian, Paul. 2014. "What Is Inference?" *Philosophical Studies* 169 (1): 1–18.

Brandom, Robert B. 1994. *Making It Explicit*. Harvard University Press.

———. 2000. *Articulating Reasons: An Introduction to Inferentialism*. Harvard University Press.

Carroll, Lewis. 1895. "What the Tortoise Said to Achilles." *Mind* 4 (14): 278–80.

Collins, P.H. 2000. *Black Feminist Thought: Knowledge, Consciousness, and the Politics of Empowerment*, 2nd ed., New York, NY: Routledge.

Dutilh Novaes, Catarina. 2015. "A Dialogical, Multi-Agent Account of the Normativity of Logic." *Dialectica* 69 (4): 587–609.

———. 2016. "Reductio Ad Absurdum from a Dialogical Perspective." *Philosophical Studies* 173 (10): 2605–28.

———. 2020. *The Dialogical Roots of Deduction: Historical, Cognitive, and Philosophical Perspectives on Reasoning*. Cambridge University Press.

Fintel, Kai von. 2008. "What Is Presupposition Accommodation, Again?" *Philosophical Perspectives* 22 (1): 137–70.

García-Carpintero, Manuel. 2016. "Accommodating Presuppositions." *Topoi* 35: 37–44.

Gelman, Susan A. 2003. *The Essential Child: Essentialism in everyday thought*, Oxford University Press.

Gelman, Susan A., Sarah-Jane Leslie, Alexandra M. Was and Christina M. Koch. 2015. "Children's interpretations of general quantifiers, specific quantifiers and generics." *Language, Cognition and Neuroscience* 30(4): 448–461. doi: 10.1080/23273798.2014.931591

Geurts, Bart. 2018. "Convention and Common Ground." *Mind & Language* 33 (2): 115–29.

Girle, Roderic A. 2016. "Proof and Dialogue in Aristotle." *Argumentation* 30: 289–316.

Hamblin, C. L. 1970. *Fallacies*. London: Methuen & Co Ltd.

Harman, Gilbert. 1986. *Change in View: Principles of Reasoning.* MIT Press.

Haslanger, Sally. 2007. "'But Mom, Crop-Tops *are* Cute!' Social Knowledge, Social Structure and Ideology Critique." *Philosophical Issues* 17 (1): 70–91. doi: 10.1111/j.1533-6077.2007.00123.x

———. 2011. "Ideology, Generics, and Common Ground." In *Feminist Metaphysics: Explorations in the Ontology of Sex, Gender and the Self*, edited by Charlotte Witt, 179–207. Dordrecht: Springer.

Hearn, Jeff. 1998. *The Violences of Men: how mern talk about and how agencies respond to men's violence to women*, London: Sage Publications.

Kukla, Rebecca, and Mark Lance. 2009. *Yo! And Lo!: The Pragmatic Topography of the Space of Reasons*. Harvard University Press.

Lance, Mark Norris, and John O'Leary-Hawthorne. 1997. *The Grammar of Meaning*. Cambridge University Press.

Lance, Mark, and W. Heath White. 2007. "Stereoscopic Vision: Persons, Freedom, and Two Spaces of Material Inference." *Philosophers' Imprint* 7 (4): 1–21.

Langton, Rae. 2015. "Accommodating Injustice—the 2015 Locke Lectures."

Leslie, Sarah-Jane. 2008. "Generics: Cognition and Acquisition." *The Philosophical Review* 117 (1): 1–47.

———. 2012. "Generics Articulate Default Generalizations." *Recherches Linguistiques de Vincennes* 41: 25–44.

———. 2017. "The Original Sin of Cognition: Fear, Prejudice, and Generalization." *Journal of Philosophy* 114 (8): 1–29.

Kirkpatrick, James Ravi. 2024. "The Acquisition of Generics." *Mind & Language*, 39:492–517. doi: 10.1111/mila.12499

Liebesman, David, and Rachel Katharine Sterken. 2021. "Generics and the Metaphysics of Kinds." *Philosophy Compass* 16 (7): e12754. https://doi.org/10.1111/phc3.12754.

MacFarlane, John. 2010. "Pragmatism and Inferentialism." In *Reading Brandom: On Making It Explicit*, edited by Bernhard Weiss and Jeremy Wanderer, 81–95. Routledge.

Medina, José. 2013. *The Epistemology of Resistance: gender, racial oppression, epistemic injustice, and resistant imaginations*. Oxford University Press.

———. 2017. "Varieties of Hermeneutical Injustice." In *The Routledge Handbook of Epistemic Injustice*, edited by Ian James Kidd, José Medina and Gaile Polhaus, Jr., 41–52. Routledge.

Murray, Sarah E., and William B. Starr. 2018. "Force and Conversational States." In *New Work on Speech Acts*, edited by Daniel Fogal, Daniel Harris, and Matthew Moss, 202–36. New York: Oxford University Press.

Neufeld, Eleonore. 202+. "Engineering Social Concepts: Labels and the Science of Categorization." This volume.

Nickel, Bernhard. 2016. *Between Logic and the World*. Oxford University Press. https://doi.org/10.1093/acprof:oso/9780199640003.001.0001.

———. 2017. "Generics." In *A Companion to the Philosophy of Language*, edited by Bob Hale, Crispin Wright, and Alexander Miller, Second, 437–62. John Wiley & Sons, Ltd. https://doi.org/10.1002/9781118972090.ch18.

Pain, Rachel. 2014. "Everyday Terrorism: Connecting Domestic Violence and Global Terrorism." *Progress in Human Geography* 38 (4): 531–50.

Pelletier, F., and N. Asher. 1997. "Generics and Defaults." In *Handbook of Logic and Language*, edited by J. van Benthem and A. ter Meulen, 1125–79. MIT Press.

Polhaus, Jr., Gaile. 2017. "Varieties of Epistemic Injustice." In *The Routledge Handbook of Epistemic Injustice*, edited by Ian James Kidd, José Medina and Gaile Polhaus, Jr., 13–26. Routledge.

Portner, Paul. 2004. "The Semantics of Imperatives Within a Theory of Clause Types." In *Proceedings of Semantics and Linguistic Theory 14*, edited by K. Watanabe and R. B. Young, 235–52. Ithaca, NY: CLC Publications.

———. 2009. *Modality*. Oxford; New York: Oxford University Press.

Restall, Greg. 2005. "Multiple Conclusions." In *Logic, Methodology and Philosophy of Science: Proceedings of the Twelfth International Congress*, edited by Petr Hájek, Luis Valdés-Villanueva, and Dag Westerståhl, 189–205. KCL Publications.

———. 2009. "Truth Values and Proof Theory." *Studia Logica* 92 (2): 241–64.

———. 2022. *Proofs and Models in Philosophical Logic*, Cambridge University Press.

Roberts, Craige. 2012. "Information Structure in Discourse: Towards an Integrated Formal Theory of Pragmatics." *Semantics and Pragmatics* 5 (6): 1–69. https://doi.org/10.3765/sp.5.6.

———. 2015. "Accommodation in a Language Game." In *A Companion to David Lewis*, edited by Barry Loewer and Jonathan Schaffer, 345–66. John Wiley & Sons, Ltd.

———. 2018. "Speech Acts in Discourse Context." In *New Work on Speech Acts*, edited by Daniel Fogal, Daniel Harris, and Matt Moss, 317–59. Oxford University Press.

Simons, Mandy. 2003. "Presupposition and Accommodation: Understanding the Stalnakerian Picture." *Philosophical Studies* 112 (3): 251–78.

———. 2006. "Foundational Issues in Presupposition." *Philosophy Compass* 1 (4): 357–72.

Stalnaker, Robert. 2002. "Common Ground." *Linguistics and Philosophy* 25: 701–21.

Sterken, Rachel Katharine. 2017. "The Meaning of Generics." *Philosophy Compass* 12 (8): e12431.

Stovall, Preston. 2019. "Characterizing Generics Are Material Inference Tickets: A Proof-Theoretic Analysis." *Inquiry*, 1–37. https://doi.org/10.1080/0020174x.2019.1580839.